

Introduction to Convolutional Neural Network

Code & Data

Vinh Dinh Nguyen - PhD in Computer Science

Quoc-Thai Nguyen - TA

Outline



Neural Network: Review and Limitations

Convolution Layer

Pooling Layer

Flatten Layer

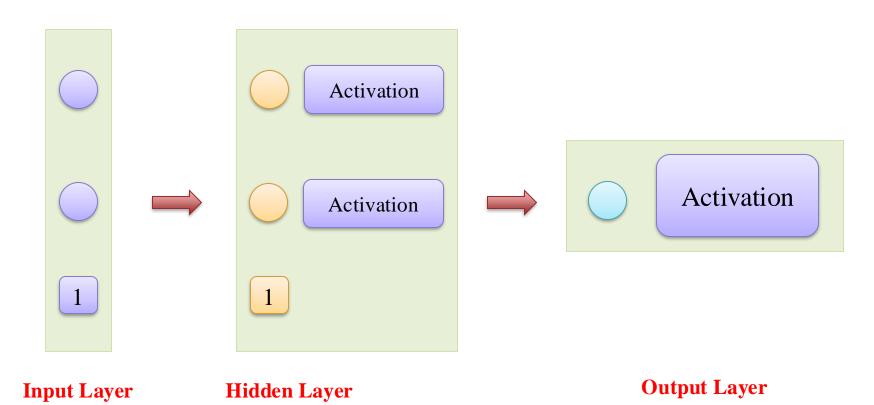
Practice

Advanced Dicussion



!

Neural Network



Loss: CrossEntropyLoss

$$L(\mathbf{\theta}) = -\sum_{i} y_{i} \log(\hat{\mathbf{y}}_{i})$$

Optimizer: SGD

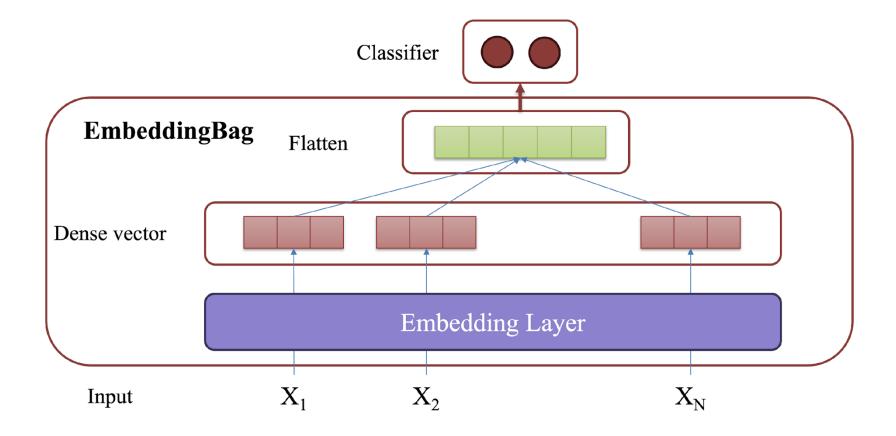
$$x = x - \eta * f'(x)$$



!

Neural Network for Text

No capture the order and importance of words in a sentence

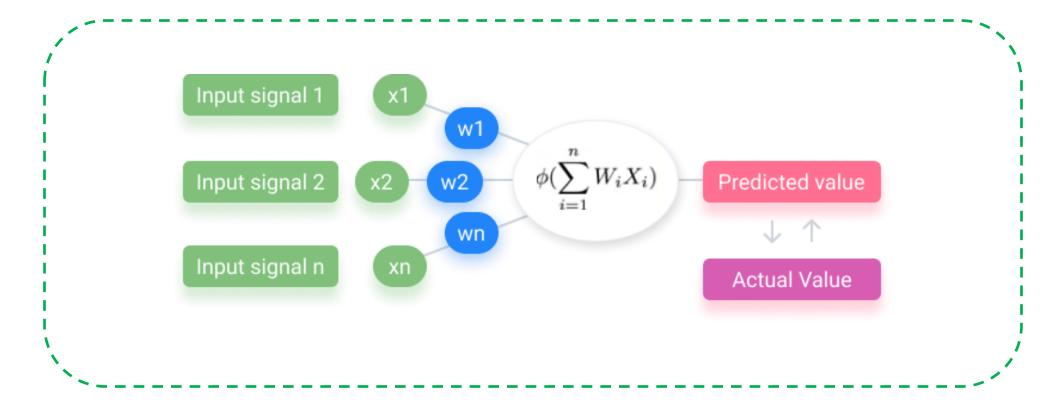




1

Neural Network for Time Series

Cannot deal with the problem of having different amount of input values

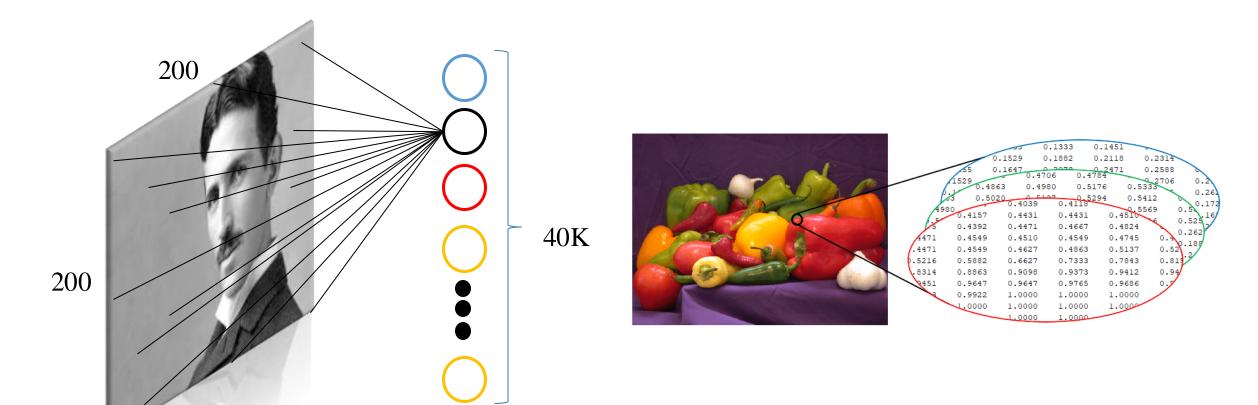




!

Neural Network for Image

Each hidden node connects to all the other nodes





1

Neural Network

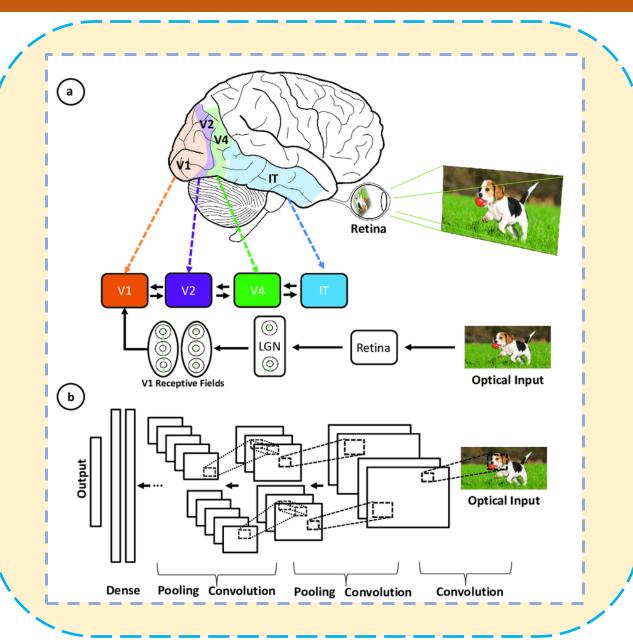
Need better network architectures...

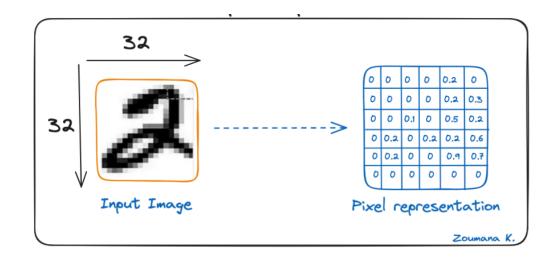
RNNs for Sequence

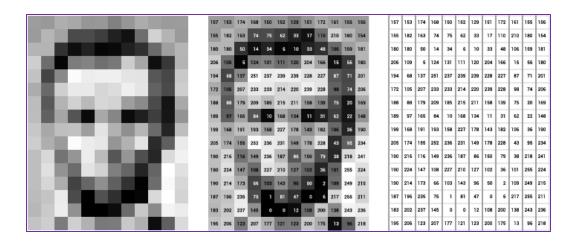
CNNs for Image



CNN Motivation







Outline

Neural Network: Review and Limitations



Convolution Layer

Pooling Layer

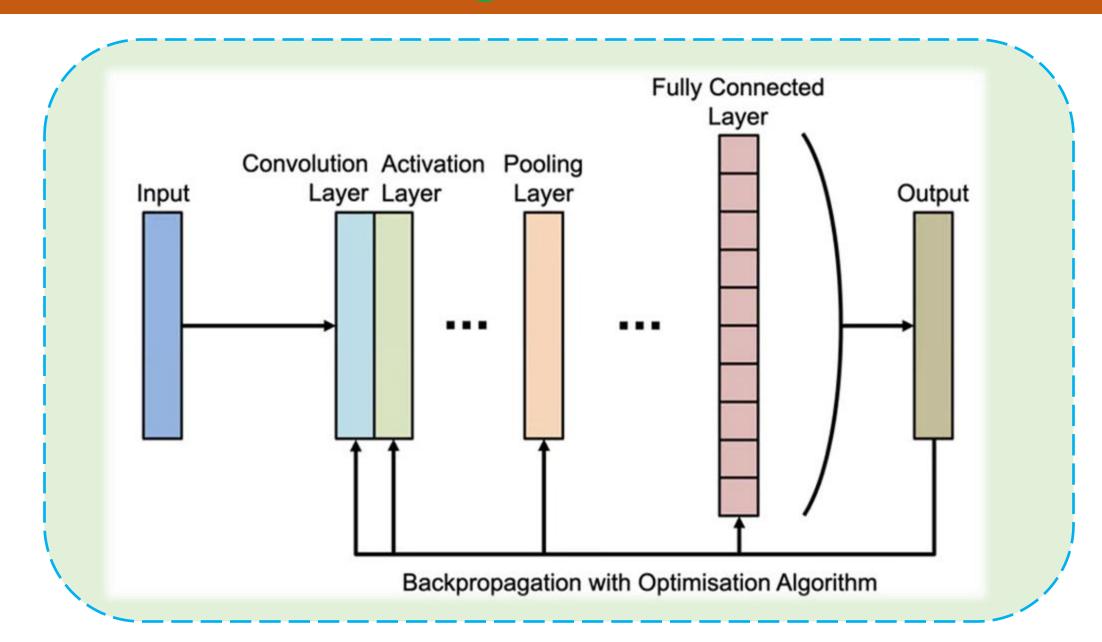
Flatten Layer

Practice

Advanced Dicussion



The Building Blocks of a CNN





Convolution Layer: Quick Look



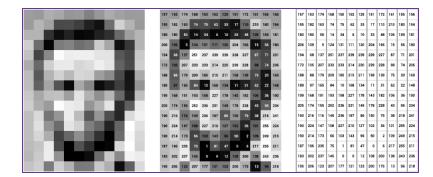
Convolution Kernel



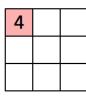


$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



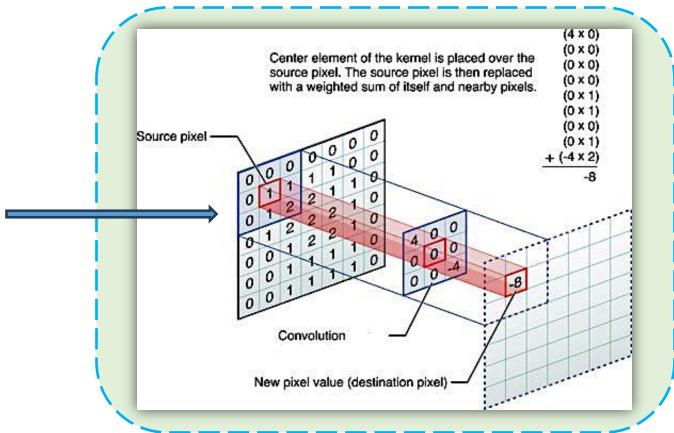


1,	1,0	1 _{×1}	0	0
0,×0	1,	1,0	1	0
0,1	0,0	1,	1	1
0	0	1	1	0
\sim	1	1	\	^



Image

Convolved Feature



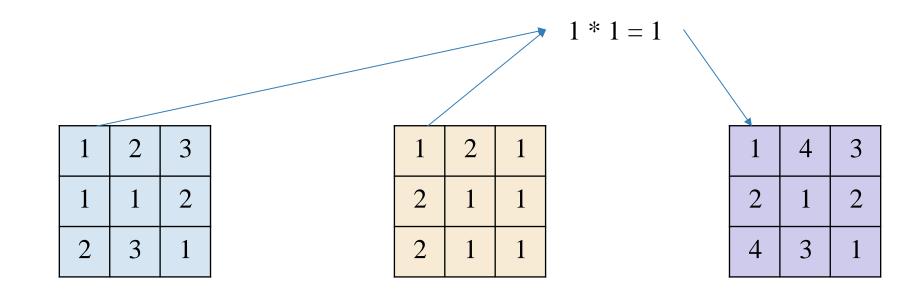
Applying filters to an image highlights its edges, creating a new image (in CNN terms, a feature map)



!

Convolutional Operation

- **Element-wise Multiplication Matrix**
 - \rightarrow A (MxN) B (MxN) => C (MxN)





!

Convolutional Operation

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

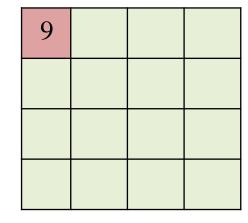
Input: 6 x 6

1	1	0
1	1	0
1	0	1

*

Kernel: 3 x 3

2x1 + 2x1 + 1x0 +
0x1 + 4x1 + 0x0 +
0x1 + 4x0 + 1x1 = 9





1

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4

Input: 6 x 6

1	1	0
1	1	0
1	0	1

*

Kernel: 3 x 3

9	13	



!

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

1	1	0
1	1	0
1	0	1

*

Kernel: 3 x 3

9	13	9	13
14	11	13	10
12	17	11	14
12	13	13	18



!

Convolutional Operation

Pytorch

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
tensor([[[2., 2., 1., 4., 1., 0.],
         [0., 4., 0., 3., 3., 4.],
         [0., 4., 1., 2., 0., 0.],
         [2., 1., 4., 1., 3., 1.],
         [4., 3., 1., 4., 2., 4.],
         [2., 0., 0., 4., 3., 4.]]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    bias=False
conv_layer.weight
Parameter containing:
tensor([[[[ 0.0520, 0.2693, 0.0364],
          [-0.1051, 0.0896, -0.0904],
          [ 0.1403, 0.2976, 0.1927]]]], requires grad=True)
```

```
init_kernel_weight = torch.randint(
    high=2,
    size=(conv_layer.weight.data.shape),
    dtype=torch.float32
init_kernel_weight
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]])
# init weight
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
Parameter containing:
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]], requires_grad=True)
output = conv_layer(input)
output
tensor([[[ 9., 13., 9., 13.],
         [14., 11., 13., 10.],
         [12., 17., 11., 14.],
         [12., 13., 13., 18.]]], grad_fn=<SqueezeBackward1>)
```



Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	1 4	3 2	4

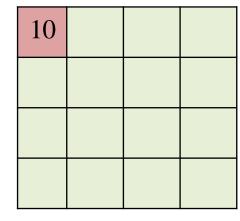
Input: 6 x 6

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

1

Bias





!

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

1

Bias

10	14	10	14
15	12	14	11
13	18	12	15
13	14	14	19



!

Convolutional Operation

Pytorch

```
input
tensor([[[2., 2., 1., 4., 1., 0.],
         [0., 4., 0., 3., 3., 4.],
         [0., 4., 1., 2., 0., 0.],
         [2., 1., 4., 1., 3., 1.],
         [4., 3., 1., 4., 2., 4.],
         [2., 0., 0., 4., 3., 4.]]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in channels=1,
    out channels=1,
    kernel_size=3, # create a kernel: 3 x 3
init_kernel_weight
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]])
# init weight
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
Parameter containing:
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]]], requires_grad=True)
```

```
conv_layer.bias
Parameter containing:
tensor([-0.1148], requires_grad=True)
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
conv_layer.bias
Parameter containing:
tensor([1.], requires_grad=True)
output = conv_layer(input)
output
tensor([[[10., 14., 10., 14.],
         [15., 12., 14., 11.],
         [13., 18., 12., 15.],
         [13., 14., 14., 19.]]], grad fn=<SqueezeBackward1>)
```



!

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

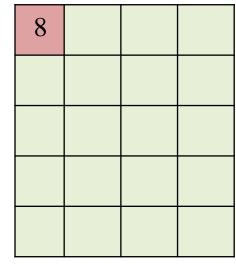
1	0	1
0	1	1

Kernel: 2 x 3

*

1

Bias





!

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

1	0	1
0	1	1

Kernel: 2 x 3

*

1

Bias

8	10	9	12
6	11	6	8
7	12	6	7
11	8	14	9
6	12	11	16





Convolutional Operation

Pytorch

```
# define convolutional layer
                                                      conv_layer.bias
conv layer = nn.Conv2d(
                                                      Parameter containing:
   in channels=1,
                                                      tensor([0.3672], requires_grad=True)
   out_channels=1,
    kernel size=(2, 3), # create a kernel: 2 x 3
                                                      # init bias
                                                      conv layer.bias = nn.Parameter(
                                                          torch.tensor([1], dtype=torch.float32)
# init weight & bias
conv_layer.weight.data = init_kernel_weight
                                                      conv layer bias
conv_layer.weight
                                                      Parameter containing:
Parameter containing:
                                                      tensor([1.], requires_grad=True)
tensor([[[[1., 0., 1.],
          [0., 1., 1.]]]], requires_grad=True)
                      output = conv_layer(input)
                      output
                      tensor([[[ 8., 10., 9., 12.],
                                [6., 11., 6., 8.],
                                [7., 12., 6., 7.],
                                [11., 8., 14., 9.],
```

[6., 12., 11., 16.]]], grad_fn=<SqueezeBackward1>)



!

Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: M x N

1	0	1
0	1	1

Kernel: K x O

*

1

Bias

8	10	9	12
6	11	6	8
7	12	6	7
11	8	14	9
6	12	11	16

Output:

$$M - (K - 1) \times N - (O - 1)$$



!

Padding

*

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input: 4 x 4

Padding: 1 x 1

0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0

Shape: 6 x 6

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1

Bias

7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6



!

Padding

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
tensor([[[2., 3., 1., 4.],
         [1., 1., 3., 2.].
         [0., 4., 3., 0.],
         [3., 2., 2., 0.]]])
init_kernel_weight = torch.randint(
    high=2,
    size=(conv layer.weight.data.shape),
    dtype=torch.float32
init_kernel_weight
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding='same'
                             {"valid", "same"}
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires grad=True)
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
conv_layer.bias
Parameter containing:
tensor([1.], requires grad=True)
output = conv_layer(input)
output
tensor([[[ 7., 8., 12., 8.],
         [8., 16., 18., 11.],
         [10., 15., 16., 9.],
         [10., 15., 12., 6.]]], grad fn=<SqueezeBackward1>)
```



!

Padding

*

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input: 4 x 4

Padding: 2 x 1

0	0	0	0	0	0
0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Shape: 8 x 6

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1

Bias

3	4	2	5
7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6
6	8	5	3



!

Padding

```
input
tensor([[[2., 3., 1., 4.],
         [1., 1., 3., 2.],
         [0., 4., 3., 0.],
         [3., 2., 2., 0.]]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding=(2, 1)
                           An int / a tuple of ints
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires_grad=True)
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
conv_layer.bias
Parameter containing:
tensor([1.], requires_grad=True)
output = conv_layer(input)
output
tensor([[[ 3., 4., 2., 5.],
         [7., 8., 12., 8.],
         [8., 16., 18., 11.],
         [10., 15., 16., 9.],
         [10., 15., 12., 6.],
         [ 6., 8., 5., 3.]]], grad fn=<SqueezeBackward1>)
```



!

Padding

*

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input: M x N

Padding: P x Q

0	0	0	0	0	0
0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Shape: $(M+2P) \times (N+2Q)$

1	1	1
1	1	1
0	1	0

Kernel: K x O

1

Bias

3	4	2	5
7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6
6	8	5	3

Output:

 $(M+2P-K+1) \times (N+2Q-O+1)$



!

Stride

Stride: 1 (1x1)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

*



Bias

10	10	13	15
10	12	14	15
11	12	16	17
12	16	14	16



!

Stride

Stride: 2 (2x2)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	$\frac{2}{0}$

Input: 6 x 6

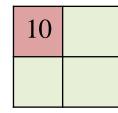
1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

*

1

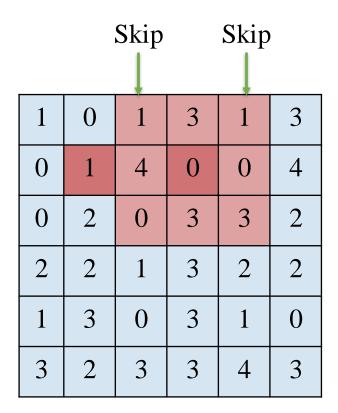
Bias





!

Stride



Input: 6 x 6

Stride: 2 (2x2)

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

*

1	1	l
---	---	---

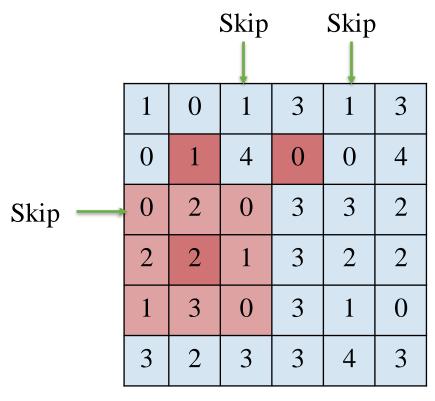
Bias

10	13





Stride



Input: 6 x 6

Stride: 2 (2x2)

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

*

	1	

Bias

10	13
11	



!

Stride

Stride: 2 (2x2)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

*

	1	

Bias

10	13
11	16



!

Stride

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
tensor([[[1., 0., 1., 3., 1., 3.],
         [0., 1., 4., 0., 0., 4.],
         [0., 2., 0., 3., 3., 2.],
         [2., 2., 1., 3., 2., 2.],
         [1., 3., 0., 3., 1., 0.],
         [3., 2., 3., 3., 4., 3.]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    stride=2
```

```
conv_layer.weight.data = init_kernel_weight
conv layer.weight
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires grad=True)
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
conv_layer.bias
Parameter containing:
tensor([1.], requires_grad=True)
output = conv_layer(input)
output
tensor([[[10., 13.],
         [11., 16.]]], grad_fn=<SqueezeBackward1>)
```



!

Stride

Stride: 2 (2x2)

*

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

Input: 4 x 4

Padding: 1 x 1

0	0	0	0	0	0
0	0	3	1	1	0
0	3	1	2	0	0
0	3	4	2	3	0
0	3	0	0	2	0
0	0	0	0	0	0

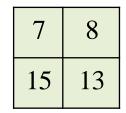
Shape: 6 x 6

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3



Bias





!

Stride

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
tensor([[[0., 3., 1., 1.],
         [3., 1., 2., 0.],
         [3., 4., 2., 3.],
         [3., 0., 0., 2.]]])
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding=1,
    stride=(2, 2)
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires grad=True)
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
conv_layer.bias
Parameter containing:
tensor([1.], requires_grad=True)
output = conv_layer(input)
output
tensor([[[ 7., 8.],
         [15., 13.]]], grad fn=<SqueezeBackward1>)
```



Convolutional Layer

!

Stride

Stride: (S, T)

*

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

Input: M x N

Padding: (P, Q)

0	0	0	0	0	0
0	0	3	1	1	0
0	3	1	2	0	0
0	3	4	2	3	0
0	3	0	0	2	0
0	0	0	0	0	0

Shape: $(M+2P) \times (N+2Q)$

1	1	1
1	1	1
0	1	0

Kernel: K x O

$$\left| \frac{M + 2P - K}{S} + 1 \right| \times \left| \frac{N + 2Q - K}{T} + 1 \right|$$

Outline

Neural Network: Review and Limitations

Convolution Layer



Pooling Layer

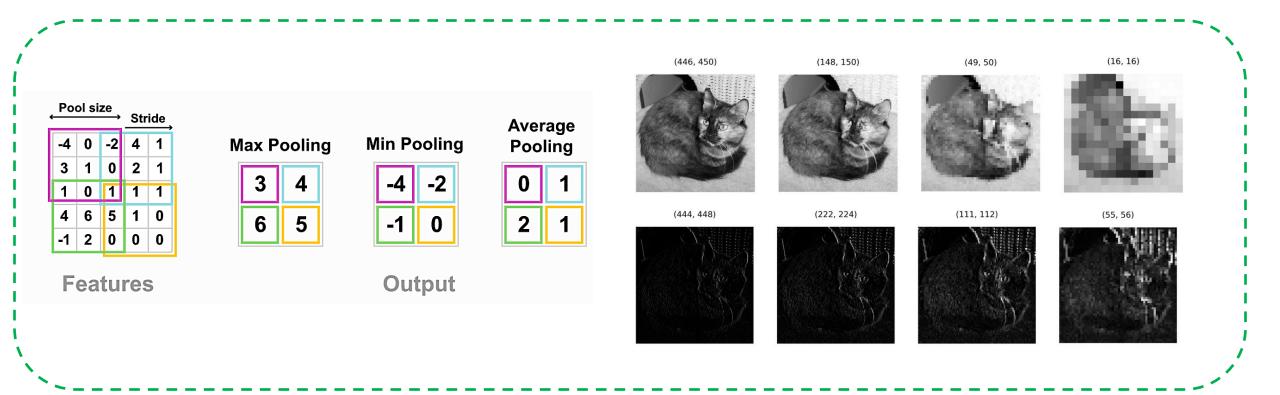
Flatten Layer

Practice

Advanced Dicussion



Pooling Layer: Quick Look



This layer helps to make the CNN more computationally efficient by reducing the number of parameters and ensuring that the model focuses on the most important features.



!

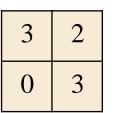
Max Pooling

Kernel Size: 2

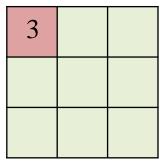
Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6



Max values





!

Max Pooling

Kernel Size: 2

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

1	0	
3	1	

Max values

3	3	



!

Max Pooling

Kernel Size: 2

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

1	0	
3	1	

Max values

3	3	3
4	4	4
4	4	4



!

Max Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2 Stride: 2

```
    3
    3

    4
    4

    4
    4
```

Output: 3 x 3

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]])
max_pool_layer = nn.MaxPool2d(kernel_size=2)
                                     Default: Stride = 2
output = max_pool_layer(input)
output
tensor([[[3., 3., 3.],
         [4., 4., 4.],
```

[4., 4., 4.]]])



!

Max Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2 Stride: (1, 2)

```
    3
    3

    3
    4

    4
    4

    4
    3

    4
    4
```

```
input
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]]])
max_pool_layer = nn.MaxPool2d(
    kernel_size=2,
    stride=(1, 2)
output = max_pool_layer(input)
output
tensor([[[3., 3., 3.],
         [3., 4., 1.],
         [4., 4., 4.],
         [4., 3., 4.],
         [4., 4., 4.]]])
```



1

Max Pooling

MaxPool1d

Kernel Size: 3

Stride: 3

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

```
      3
      3

      3
      1

      4
      1

      4
      4

      3
      3

      4
      4
```

Output: 6 x 2

```
input
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]])
max_pool_layer = nn.MaxPool1d(
    kernel_size=3,
    stride=3
max_pool_layer(input)
tensor([[[3., 3.],
         [3., 1.],
         [4., 1.],
         [4., 4.],
         [3., 3.],
```

[4., 4.]]])



!

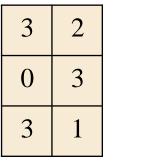
Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6



Average values

2.0	



!

Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

1	0	
3	1	
4	1	

Average values

2	1.7	



!

Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

3	1	
2	4	
1	0	

Average values

2	1.7	0.8
1.8		



1

Average Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: (3, 2) Stride: 2

2	1.7	0.8
1.8	1.6	1.3

```
input
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]])
avg_pool_layer = nn.AvgPool2d(
    kernel_size=(3, 2),
    stride=(2, 2)
output = avg_pool_layer(input)
output
tensor([[[2.0000, 1.6667, 0.8333],
         [1.8333, 1.6667, 1.3333]]])
```



!

Average Pooling

AvgPool1d

Kernel Size: 3

Stride: 3

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

```
    2.0
    1.0

    2.0
    0.7

    2.7
    0.7

    2.3
    1.7

    1.3
    1.0

    3.7
    3.3
```

```
input
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]])
avg_pool_layer = nn.AvgPool1d(
    kernel_size=3,
    stride=3
output = avg_pool_layer(input)
output
tensor([[[2.0000, 1.0000],
         [2.0000, 0.6667],
         [2.6667, 0.6667],
         [2.3333, 1.6667],
         [1.3333, 1.0000],
         [3.6667, 3.3333]]])
```

Outline

Neural Network: Review and Limitations

Convolution Layer

Pooling Layer



Flatten Layer

Practice

Advanced Dicussion



Flatten Layer

!

Flattens a contiguous range of dims into a tensor

Input: 3 x 2

```
2 4 3 1 3 4
```

Output: 1 x 6

input = torch.randint(5, (1, 3, 2), dtype=torch.float32)

Outline

Neural Network: Review and Limitations

Convolution Layer

Pooling Layer

Flatten Layer



Practice

Advanced Dicussion



1

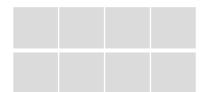
Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3





1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

4 7



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

4 7 5



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

4 7 5 8



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

4	7	5	8
4	8	4	8



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

*

2



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

*

6	9	7	10
6	10	6	10

2



2

Exercise – Padding

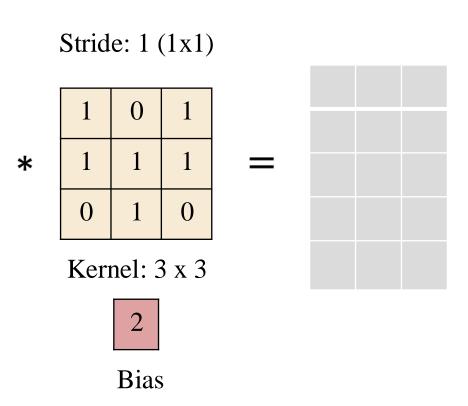
2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 4 x 6

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 6 x 8





2

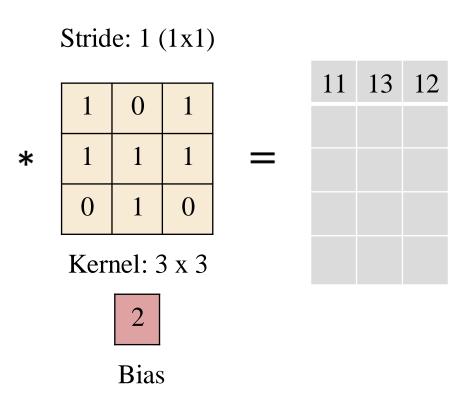
Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0



2

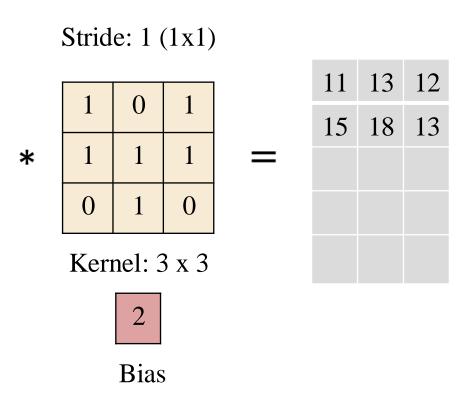
Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0





2

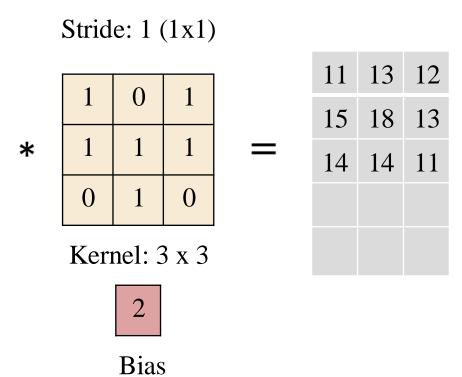
Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0





2

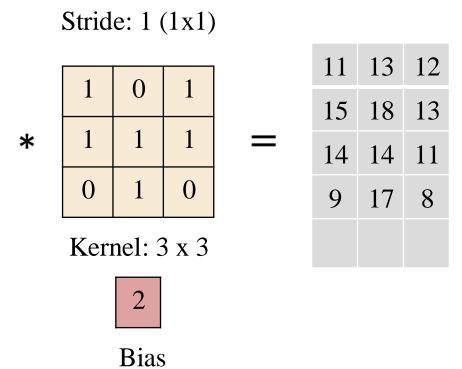
Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0



2

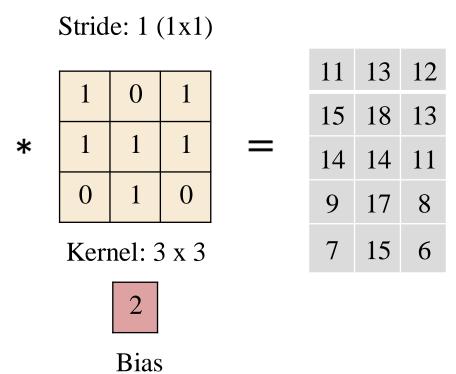
Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0





2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)

	1	0	1		
k	1	1	1	=	
	0	1	0		

Kernel: 3 x 3





2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)

1	0	1		11	12
1	1	1	=		
0	1	0			

Kernel: 3 x 3

*

2

2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

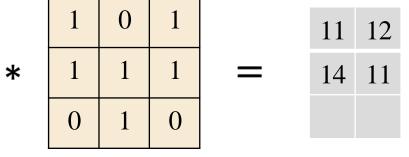
Padding: 1 x 1

Input: 5 x 3

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)



Kernel: 3 x 3

2

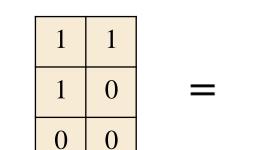
2

Exercise – Convolutional Layer + Pooling

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

*

Input: 5 x 3



Kernel: 3 x 2

Stride: 1 (1x1)





2

Exercise – Convolutional Layer + Pooling

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

*

Input: 5 x 3

Stride: 1 (1x1)

1	1	
1	0	=
0	0	

Kernel: 3 x 2



Bias

Max Pooling Kernel Size: (1x2)





2

Exercise – Convolutional Layer + Pooling

8

8

6

10

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

*

Input: 5 x 3

Stride: 1 (1x1)

1	1	
1	0	=
0	0	

Kernel: 3 x 2



Bias

Max Pooling Kernel Size: (1x2)

10

8

6

2

Exercise – Pooling For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

MaxPooling =

Output: 3 x 3

Input: 7 x 7



2

Exercise – Pooling For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

MaxPooling = 2x2

0	43	43
191	250	251
241	198	250

Output: 3 x 3

Input: 7 x 7



2

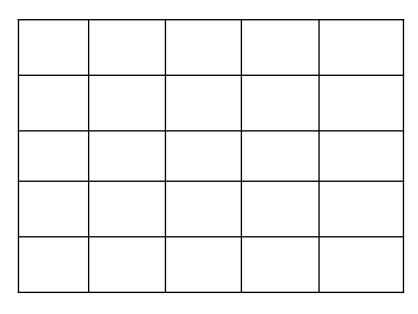
Exercise – Pooling For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

*

1	0	-1
1	0	-1
1	0	-1

Kernel: 3 x 3



Output: 5 x 5

Input: 7 x 7



2

Exercise – Pooling For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

*

1	0	-1
1	0	-1
1	0	-1

Kernel: 3 x 3

-250	-243	82	22	168
-288	34	206	-59	168
212	657	294	185	244
-155	248	29	64	202
-193	127	229	486	202

Output: 5 x 5



2

Exercise – Pooling For Grayscale Image

-250	-243	82	22	168
-288	34	206	-59	168
212	657	294	185	244
-155	248	29	64	202
-193	127	229	486	202

MaxPooling Kernel: 2

34	206
657	297

Input: 5 x 5

Outline

Neural Network: Review and Limitations

Convolution Layer

Pooling Layer

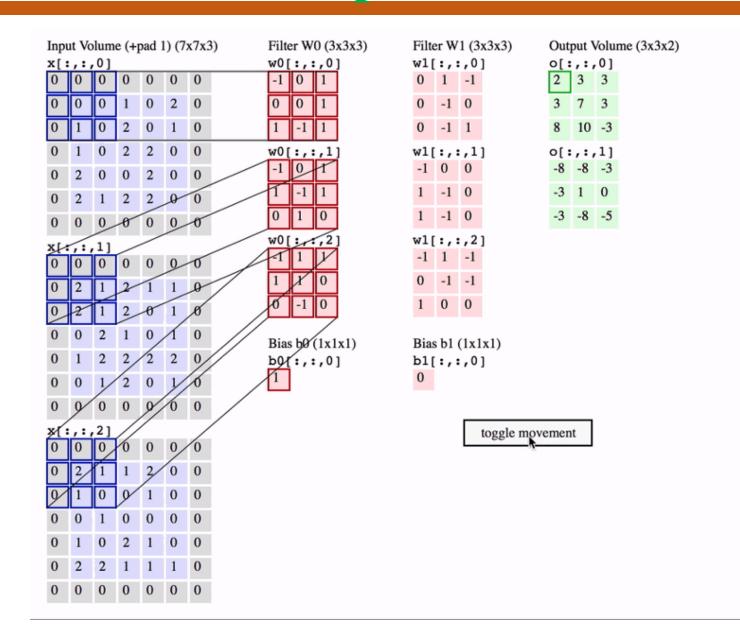
Flatten Layer

Practice



Advanced Dicussion







SL.No		Activation Shape	Activation Size	# Parameters
1.	Input Layer:	(32, 32, 3)	3072	
2.	CONV1 (f=5, s=1)	(28, 28, 8)	6272	
3.	POOL1	(14, 14, 8)	1568	
4.	CONV2 (f=5, s=1)	(10, 10, 16)	1600	
5.	POOL2	(5, 5, 16)	400	
6.	FC3	(120, 1)	120	
7.	FC4	(84, 1)	84	
8.	Softmax	(10, 1)	10	



